

EDA

August 10, 2022

0.1 EDA Performed on the Data

```
[3]: # Read the csv dataset
data = pd.read_csv('Healthcare_dataset.csv')

# Drop the ID variable
data = data.drop(["PtId"], axis=1)

data.head()
```

```
[3]:  Persistency_Flag  Gender      Race  Ethnicity  Region  Age_Bucket  \
0      Persistent   Male    Caucasian  Not Hispanic    West      >75
1  Non-Persistent   Male      Asian  Not Hispanic    West    55-65
2  Non-Persistent  Female  Other/Unknown    Hispanic  Midwest    65-75
3  Non-Persistent  Female    Caucasian  Not Hispanic  Midwest    >75
4  Non-Persistent  Female    Caucasian  Not Hispanic  Midwest    >75
```

```
      Ntm_Speciality  Ntm_Specialist_Flag      Ntm_Speciality_Bucket  \
0  GENERAL PRACTITIONER      Others  OB/GYN/Others/PCP/Unknown
1  GENERAL PRACTITIONER      Others  OB/GYN/Others/PCP/Unknown
2  GENERAL PRACTITIONER      Others  OB/GYN/Others/PCP/Unknown
3  GENERAL PRACTITIONER      Others  OB/GYN/Others/PCP/Unknown
4  GENERAL PRACTITIONER      Others  OB/GYN/Others/PCP/Unknown
```

```
      Gluco_Record_Prior_Ntm  ...  Risk_Family_History_Of_Osteoporosis  \
0                          N  ...                                      N
1                          N  ...                                      N
2                          N  ...                                      N
3                          N  ...                                      N
4                          Y  ...                                      N
```

```
      Risk_Low_Calcium_Intake  Risk_Vitamin_D_Insufficiency  \
0                          N                          N
1                          N                          N
2                          Y                          N
3                          N                          N
4                          N                          N
```

	Risk_Poor_Health_Frailty	Risk_Excessive_Thinness	\
0	N	N	
1	N	N	
2	N	N	
3	N	N	
4	N	N	

	Risk_Hysterectomy_Oophorectomy	Risk_Estrogen_Deficiency	Risk_Immobilization	\
0	N	N	N	
1	N	N	N	
2	N	N	N	
3	N	N	N	
4	N	N	N	

	Risk_Recurring_Falls	Count_Of_Risks
0	N	0
1	N	0
2	N	2
3	N	1
4	N	1

[5 rows x 68 columns]

1 Data Cleaning

```
[4]: # Total number of missing values
data.isnull().sum().sum()
```

[4]: 0

Note: The data does not contain any missing values.

```
[5]: # Dimension of the dataset
data.shape
```

[5]: (3424, 68)

```
[6]: # Standardize column names

columns = list(data.columns)

for item in columns:
    special_characters = "!@#$$%^&*()-+?=,<>/"

    for character in special_characters:
        if character in item:
            print(f"Feature: {item}, is not well formatted")
```

Feature: Comorb_Encntr_For_General_Exam_W_0_Complaint,_Susp_Or_Reprtd_Dx, is not well formatted

```
[7]: # Remove the "," from this feature name

data["Comorb_Encntr_For_General_Exam_W_0_Complaint_Susp_Or_Reprtd_Dx"] =
    ↳data["Comorb_Encntr_For_General_Exam_W_0_Complaint,_Susp_Or_Reprtd_Dx"]
data = data.
    ↳drop(["Comorb_Encntr_For_General_Exam_W_0_Complaint,_Susp_Or_Reprtd_Dx"],
    ↳axis=1)
data.shape
```

```
[7]: (3424, 68)
```

```
[8]: # Numeric Columns
numeric_col = list(data._get_numeric_data().columns)
numeric_col
```

```
[8]: ['Dexa_Freq_During_Rx', 'Count_Of_Risks']
```

1.0.1 Remove Outliers

```
[9]: data['Dexa_Freq_During_Rx'].value_counts()
```


```
[9]: 0      2488
      5      114
      6      107
      7       93
      8       71
      4       68
     10       55
     12       52
      3       46
     14       38
      9       32
     11       30
      1       24
      2       24
     13       19
     20       15
     16       14
     18       14
     22       13
     26       10
     24       10
     15        9
     30        7
     17        7
```

28	7
21	7
36	5
19	3
42	3
32	3
34	3
52	2
48	2
58	2
25	2
39	2
88	2
54	1
146	1
50	1
35	1
44	1
108	1
72	1
40	1
68	1
45	1
38	1
69	1
118	1
66	1
110	1
33	1
23	1
27	1
81	1
37	1
29	1

Name: Dexa_Freq_During_Rx, dtype: int64

```
[10]: # Spot Outliers for Dexa_Freq_During_Rx
sns.distplot(data['Dexa_Freq_During_Rx'])
```

```
[10]: <AxesSubplot:xlabel='Dexa_Freq_During_Rx', ylabel='Density'>
```



Note: The outliers in the data are skewed towards the right, to fix this we remove the 97th percentile.

```
[11]: # To remove the 99th percentile
q = data['Dexa_Freq_During_Rx'].quantile(0.97)
data_1 = data[data['Dexa_Freq_During_Rx'] < q]
data_1['Dexa_Freq_During_Rx'].describe()
```

```
[11]: count    3308.000000
      mean      1.895707
      std       3.835797
      min       0.000000
      25%       0.000000
      50%       0.000000
      75%       0.000000
      max       19.000000
      Name: Dexa_Freq_During_Rx, dtype: float64
```

```
[12]: data_1['Dexa_Freq_During_Rx'].value_counts()
```

```
[12]: 0      2488
      5      114
      6      107
      7       93
      8       71
      4       68
```

```


10      55
12      52
3       46
14      38
9       32
11      30
1       24
2       24
13      19
18      14
16      14
15       9
17       7
19       3

```

Name: Dexa_Freq_During_Rx, dtype: int64

```
[13]: sns.distplot(data_1['Dexa_Freq_During_Rx'])
```

```
[13]: <AxesSubplot:xlabel='Dexa_Freq_During_Rx', ylabel='Density'>
```



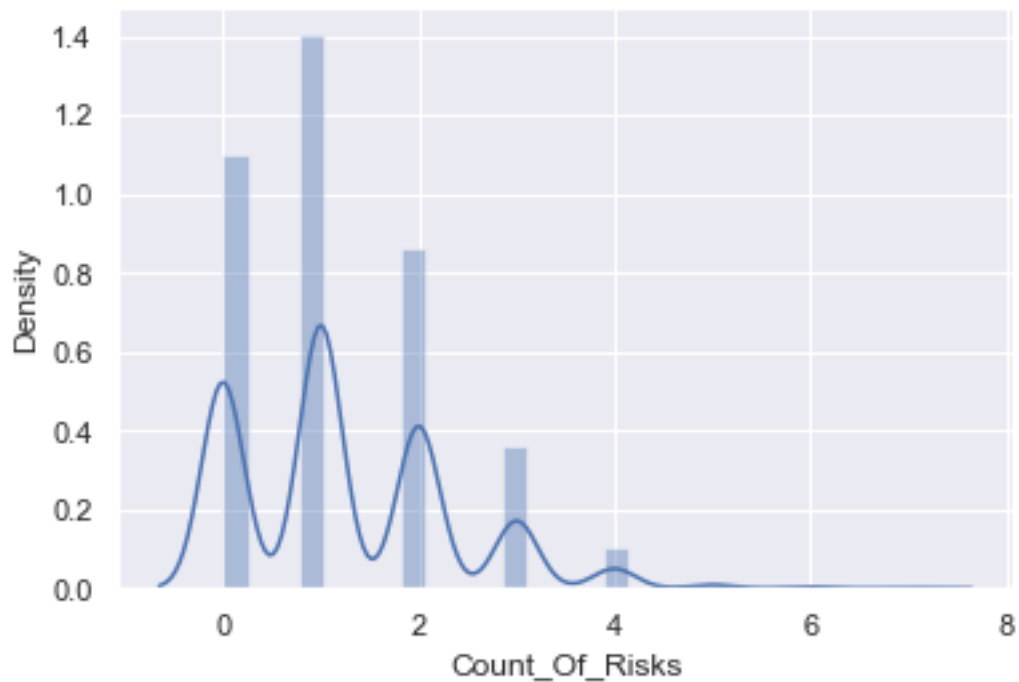
Note: The variable now contains less outliers.

```
[14]: data_1['Count_Of_Risks'].value_counts()
```

```
[14]: 1    1203
      0    943
      2    742
      3    308
      4     89
      5     15
      6      6
      7      2
      Name: Count_Of_Risks, dtype: int64
```

```
[15]: # Spot Outliers for Count_Of_Risks
      sns.distplot(data_1['Count_Of_Risks'])
```

```
[15]: <AxesSubplot:xlabel='Count_Of_Risks', ylabel='Density'>
```



Note: In this variable, the outliers are also skewed towards the right. To fix this we remove the 99th percentile.

```
[16]: # To remove the 99th percentile
      q = data_1['Count_Of_Risks'].quantile(0.99)
      data_2 = data_1[data_1['Count_Of_Risks'] < q]
      data_2.describe()
```

```
[16]:      Dexa_Freq_During_Rx  Count_Of_Risks
count              3196.000000           3196.000000
```


mean	1.875782	1.129850
std	3.834596	0.946638
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	1.000000
75%	0.000000	2.000000
max	19.000000	3.000000

```
[17]: data_2['Count_Of_Risks'].value_counts()
```

```
[17]: 1    1203
      0    943
      2    742
      3    308
      Name: Count_Of_Risks, dtype: int64
```

```
[18]: sns.distplot(data_2['Count_Of_Risks'])
```

```
[18]: <AxesSubplot:xlabel='Count_Of_Risks', ylabel='Density'>
```



This variable no longer contains outliers.

```
[19]: # Data without outliers
      data = data_2
```



```
[20]: # Data containing only categorical variables
categoric_data = data.drop(numeric_col, axis=1)
```

```
[21]: # Categorical Columns
cat_columns = list(categoric_data.columns)
```

1.0.2 Encode Target Variable

```
[23]: # The target variable
data["Persistency_Flag"].unique()
```

```
[23]: array(['Persistent', 'Non-Persistent'], dtype=object)
```

```
[24]: # Encode the Target variable

# data["Persistency_Flag"] = data["Persistency_Flag"].map({"Non-Persistent":0,
↪ "Persistent":1})

from sklearn.preprocessing import LabelEncoder

lb_make = LabelEncoder()
data["Persistency_Flag"] = lb_make.fit_transform(data["Persistency_Flag"])

# data["Persistency_Flag"].head()

data["Persistency_Flag"].unique()
```

<ipython-input-24-f18f0a11ba41>:8: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
data["Persistency_Flag"] = lb_make.fit_transform(data["Persistency_Flag"])

```
[24]: array([1, 0])
```

Note: 1 = Persistent, 0 = Non-Persistent

```
[25]: data["Persistency_Flag"].unique()
```

```
[25]: array([1, 0])
```

1.0.3 Balance Dataset

```
[26]: data["Persistency_Flag"].value_counts()
```

```
[26]: 0    2070
      1    1206
      Name: Persistency_Flag, dtype: int64
```

Note: The Non-Persistent observations are almost double the Persistent observations.

Next step is to balance the data using either the Random-Undersampling or Random_Oversampling method.

```
[27]: # Oversampling

      # import library
      from imblearn.over_sampling import RandomOverSampler

      ros = RandomOverSampler(random_state=42)

      # fit predictor and target variable
      x_ros, y_ros = ros.fit_resample(x, y)
      x_temp, y_temp = ros.fit_resample(data.drop(["Persistency_Flag"], axis=1),
      ↪data["Persistency_Flag"])
```

```
[28]: data = pd.concat([x_temp, y_temp], axis=1)
```

```
[29]: data["Persistency_Flag"].value_counts()
```

```
[29]: 1    2070
      0    2070
      Name: Persistency_Flag, dtype: int64
```

The data is now balanced, with 2070 observations of both classes each.